

Dimension reduction

Emma Yu

Feb 8, 2017

Austin ACM SIGKDD

Advanced Machine Learning with Python (Session 13)

Today's slides and the notebook:
[https://github.com/ExLupi/](https://github.com/ExLupi/ACM_SIGKDD_dimension_reduction)
[ACM_SIGKDD_dimension_reduction](https://github.com/ExLupi/ACM_SIGKDD_dimension_reduction)

Outline



- Why dimension reduction?
- Principal Component Analysis (PCA)
- PCA demo with the MNIST dataset
- Other techniques — Probabilistic PCA, LDA, Kernel PCA and Autoencoders

What are the problems we face dealing with high-dimensional data?

Problems dealing with high-dimensional data



- Expensive to store and expensive to compute (e.g. larger matrixes to inverse)
- Easier to overfit the data (too many degrees of freedom)
- Hard to visualize (we live in a 3D world)
- The curse of dimensionality — it's hard to find structures in high-dimensional data

The Curse of Dimensionality



The Curse of Dimensionality

Distance-based algorithms (e.g. KNN) do not work well in high dimensional space



The actual data usually lies near a low-dimension manifold

- We **moved and rotated** a 100*100 pixel image of 3 to make this dataset. Each image is represented by a point in the **10,000-dimensional** data space. What is the actual dimension of data?




Image credit: Bishop 2006

The actual data usually lies near a low-dimension manifold

- We **moved and rotated** a 100*100 pixel image of 3 to make this dataset. Each image is represented by a point in the **10,000-dimensional** data space. What is the actual dimension of data?

3 - x, y and rotation




Image credit: Bishop 2006

Dimension reduction:

**Deriving a set of degrees of freedom which
can be used to best re-express the dataset**

Is there another basis, which is a **linear combination** of the original basis, that preserves the **variance** of our dataset?

Is there another basis, which is a **linear combination** of the original basis, that preserves the **variance** of our dataset?

Yes! It's made of principal components!

Principal Component Analysis




Image credit: Bishop 2006

- maximizes the variance of projected data
- minimizes the average projection cost (mean squared reconstruction error)

- Also known as KLT transform, empirical orthogonal functions, quasi harmonic modes (Wikipedia)

How to find the principal components?

- Define a unit vector in the D-dimensional space \mathbf{u}_1 .
- Projected value of a data point onto this direction is $\mathbf{u}_1^T \mathbf{x}_n$.
Variance of projected data is $\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$, where S is the covariance matrix
$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T.$$
- Goal: to maximize $\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$ under the constraint $\mathbf{u}_1^T \mathbf{u}_1 = 1$.
- Solve the optimization problem, and we get:

$$\mathbf{S} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1$$

Eigenvalues and Eigenvectors

$$A\mathbf{x} = \lambda\mathbf{x}$$

Eigenvalue
↓
↑ ↑
Eigenvector

$$A = \begin{bmatrix} 2 & 0 \\ 0 & -1 \end{bmatrix} \quad \mathbf{x}_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \mathbf{x}_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$A\mathbf{x}_1 = \begin{bmatrix} 2 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 2 \\ 0 \end{bmatrix} = 2 \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 2\mathbf{x}_1$$

Eigenvalue
↓
↑
Eigenvector

$$\begin{bmatrix} | & | & | \\ \hline | & | & | \\ \hline | & | & | \\ \hline | & | & | \end{bmatrix} = \begin{bmatrix} | & | & | \\ \mathbf{v}_1 & \mathbf{v}_2 & \mathbf{v}_3 \\ | & | & | \end{bmatrix} \begin{bmatrix} \lambda_1 & 0 & 0 \\ 0 & \lambda_2 & 0 \\ 0 & 0 & \lambda_3 \end{bmatrix} \begin{bmatrix} | & | & | \\ \mathbf{v}_1 & \mathbf{v}_2 & \mathbf{v}_3 \\ | & | & | \end{bmatrix}^{-1}$$

Image credit: datinker.com

- Principal components are the largest eigenvectors of the covariance matrix
- Value of distortion = sum of all the other eigenvalues
- Cost of full eigenvector decomposition - $O(D^3)$

- PCA assumes the data is centered at zero and is sensitive to relative scaling - one wants to standardize the data first!
- PCA prefers multiple medium sized errors than one large error => sensitive to outliers
- It works great for unlabeled data. For labeled data LDA sometimes works better
- The projection to the two/three largest principal components is commonly used for visualization

Demo time!

PCA and SVD

- Singular value decomposition (SVD) calculates the **singular vectors** of the **data matrix**
- PCA finds the principal components by solving for the **eigenvectors** of the **covariance matrix**
- The covariance matrix is simply $\mathbf{X}^T \mathbf{X} / (n - 1)$
- PCA is usually solved by SVD to avoid calculating the covariance matrix

PCA and SVD

- SVD factors a matrix into $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^*$
- \mathbf{U} and \mathbf{V}^* are two unitary matrixes and $\mathbf{\Sigma}$ is a diagonal matrix
- The right singular vectors \mathbf{V} are simply the principal directions (eigenvectors of the covariance matrix)
- The diagonal terms in $\mathbf{\Sigma}$ are related to the eigenvalues of the covariance matrix via

$$\lambda_i = s_i^2 / (n - 1)$$

Probabilistic PCA

- PCA is also the maximum likelihood solution of a probabilistic latent variable model




Image credit: Bishop 2006

- Combined with EM algorithm, probabilistic PCA allows us to deal with missing data

Linear Discriminant Analysis (LDA, or Fisher's linear discriminant)

- Basic idea: find a basis that maximizes (between class variance)/(within class variance)
- Similar to PCA: LDA uses a linear combination of features. All classes are assumed to be normally distributed and have equal class covariance.
- Unlike PCA: LDA models the differences between classes of data

PCA vs LDA




Image credit: Liam Williams's Master thesis

**What if the data is not linearly
separable?**

Kernel PCA

- Use the **Kernel trick** to construct nonlinear mappings that maximize the variance in the data

- Instead of calculating the covariance matrix $\frac{1}{m} \sum_{i=1}^m \mathbf{x}_i \mathbf{x}_i^\top$
Calculate **the covariance of the transformed data**

$$\frac{1}{m} \sum_{i=1}^m \Phi(\mathbf{x}_i) \Phi(\mathbf{x}_i)^\top$$


- After some algebra, it turns out one only needs to solve the eigenvalue problem for **the kernel function**

$$\mathbf{K} \mathbf{a}_i = \lambda_i N \mathbf{a}_i \quad k(\mathbf{x}_n, \mathbf{x}_m) = \phi(\mathbf{x}_n)^\top \phi(\tilde{\mathbf{x}}_m)$$

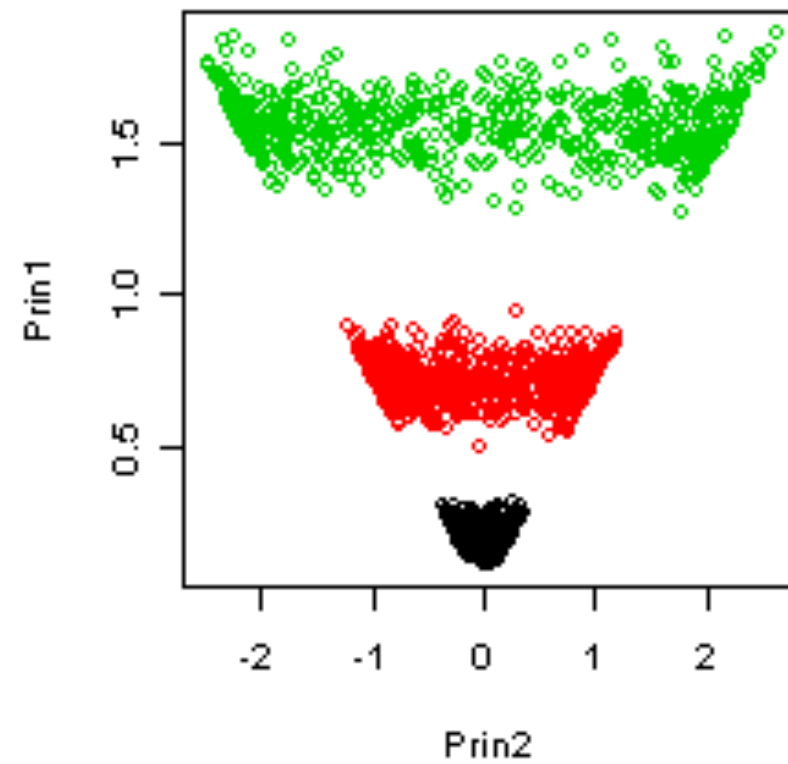
- The projection onto the eigenvector is simply

$$y_i(\mathbf{x}) = \sum_{n=1}^N a_{in} k(\mathbf{x}, \mathbf{x}_n)$$


Original



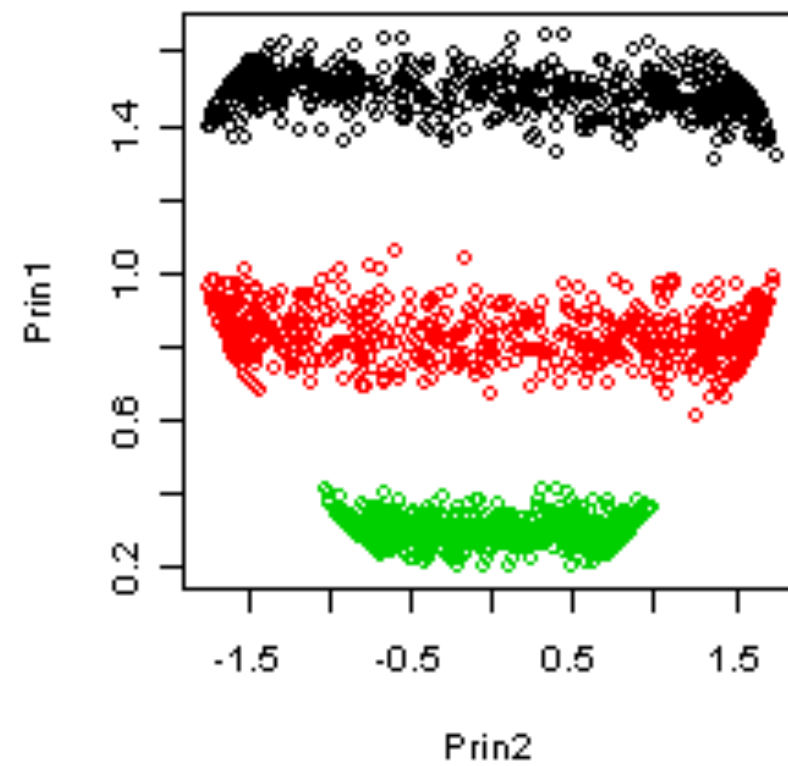
Polynomial Kernel

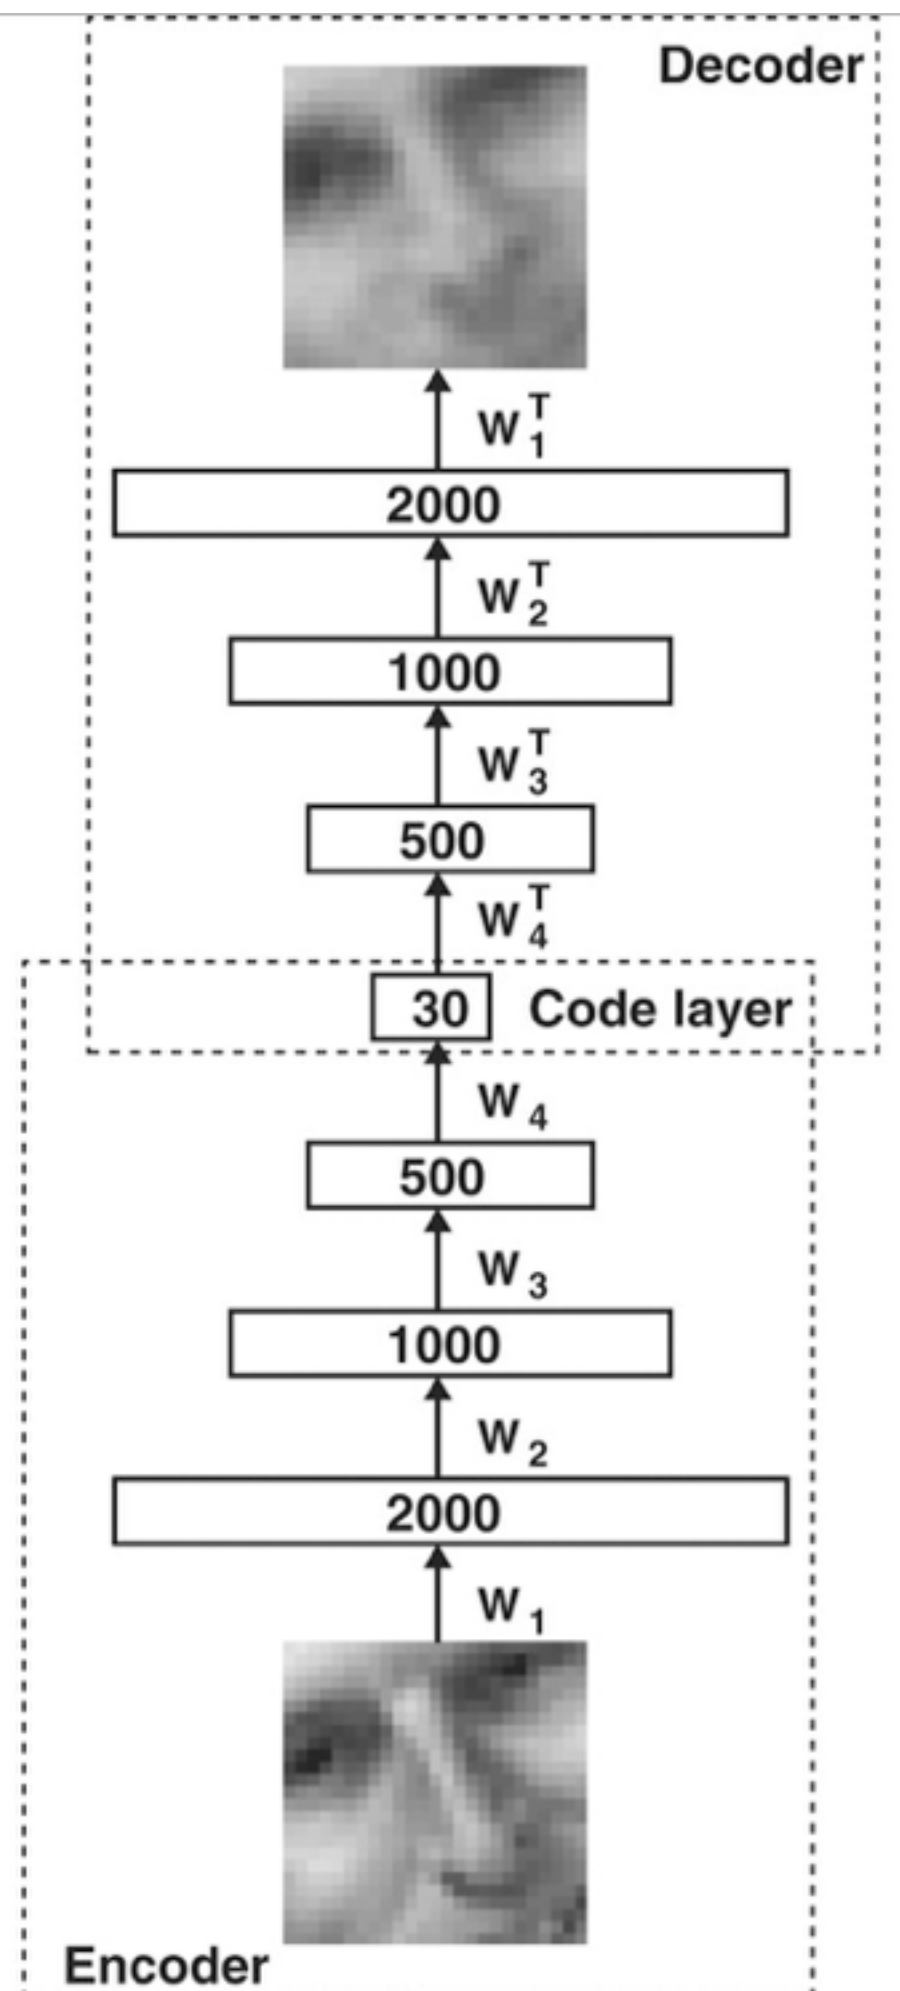


Linear PCA



Gauss Kernal





Autoencoders

- Non-linear activation function and multiple layers!
- Train to minimize reconstruction errors
- Pre-training layer-by-layer seems to be the key



Top to bottom: original data, reconstructions from a 30-D auto encoder, 30-D logistic PCA, and 30-D standard PCA

Reference

- Bishop 2006, Pattern recognition and machine leaning
- Hinton and Salakhutdinov, Science, 2006, Reducing the dimensionality of data with neural networks
- Wikipedia, Quora and Stackexchange

Contact me at:
emmayuastro@gmail.com
emmayu.com